
Multicore-Manycore Programming

Architectures, Compilers and OS/Runtimes

Fabrice Rastello, INRIA Grenoble Rhône-Alpes
Jean-François Méhaut, UJF-CEA/LIG, Grenoble

Agenda

- **Multi-Cores Architectures (JFM)**
 - Introduction
 - Memore Hierarchy
 - Instruction Level Parallelism
 - Simultaneous Multithreading (SMT)
- **Compilation, SSA (FR)**
 - Introduction to SSA
 - Properties and flavors
 - Register Allocations
 - SSA extension
- **Operating Systems for Multi-Cores (JFM)**
 - Multithreading (POSIX Threads, OpenMP,...)
 - Synchronization (Locks/semaphores, Transactional Memory,...)
- **Runtimes and multi-cores (JFM)**


What is Parallel Computing?

- **Parallel computing:** using multiple processors/cores in parallel to solve problems more quickly than with a single processor/core
- Examples of parallel machines:
 - A cluster computer that contains multiple PCs combined together with a high speed network
 - A shared memory multiprocessor (SMP*) by connecting multiple processors to a single memory system
 - A Chip Multi-Processor (CMP) contains multiple processors (called cores) on a single chip
- Concurrent execution comes from desire for performance; unlike the inherent concurrency in a multi-user distributed system
- * Technically, SMP stands for “Symmetric Multi-Processor”


Why Parallel Computing Now?

- Researchers have been using parallel computing for decades:
 - Mostly used in computational science and engineering
 - Problems too large to solve on one computer; use 100s or 1000s
- Many companies in the 80s/90s “bet” on parallel computing and failed
 - Computers got faster too quickly for there to be a large market
- Why are Universities adding courses?
 - Because the entire computing industry has bet on parallelism
 - There is a desperate need for parallel programmers
- Let’s see why...

Technology Trends: Microprocessor Capacity



2X transistors/Chip Every 1.5 years
Called "Moore's Law"




Gordon Moore (co-founder of Intel) predicted in 1965 that the transistor density of semiconductor chips would double roughly every 18 months.

Microprocessors have become smaller, denser, and more powerful.

Limit #1: Power density

Can soon put more transistors on a chip than can afford to turn on.

-- Patterson '07



Parallelism Saves Power

- Exploit explicit parallelism for reducing power

$$\text{Power} = (C * V^2 * F)/4$$

$$\text{Performance} = (\text{Cores} * F) * 1$$

Capacitance Voltage Frequency

- **Using additional cores**

- Increase density (= more transistors = more capacitance)
- Can increase cores (2x) and performance (2x)
- Or increase cores (2x), but decrease frequency (1/2): same performance at $\frac{1}{4}$ the power


- **Additional benefits**

- Small/simple cores → more predictable performance

Why Multicore Processors?


- █ Performance
- █ Consommation

1.0x





Fréquence nominale

Over-clocking



Under-clocking





Limit #2: Hidden Parallelism Tapped Out

Application performance was increasing by 52% per year as measured by the SpecInt benchmarks here

From Hennessy and Patterson,
Computer Architecture: A Quantitative Approach, 4th edition, 2006

- $\frac{1}{2}$ due to transistor density
- $\frac{1}{2}$ due to architecture changes, e.g., Instruction Level Parallelism (ILP)

VAX-11/780

- VAX : 25%/year 1978 to 1986
- RISC + x86: 52%/year 1986 to 2002

Limit #2: Hidden Parallelism Tapped Out

- Superscalar (SS) designs were the state of the art; many forms of parallelism not visible to programmer
 - multiple instruction issue
 - dynamic scheduling: hardware discovers parallelism between instructions
 - speculative execution: look past predicted branches
 - non-blocking caches: multiple outstanding memory ops
- You may have heard of these in your architecture courses, but you haven't needed to know about them to write software (C/C++, Java,...)
- Unfortunately, these sources have been used up

Performance Comparison




Figure 4. IPC Breakdown for a single 2-issue





Figure 5. IPC Breakdown for the 6-issue processor.

- Measure of success for hidden parallelism is Instructions Per Cycle (IPC)
- The 6-issue has higher IPC than 2-issue, but far less than 3x
- Reasons are: waiting for memory (D and I-cache stalls) and dependencies (pipeline stalls)


Uniprocessor Performance (SPECint) Today



- VAX : 25%/year 1978 to 1986
- RISC + x86: 52%/year 1986 to 2002
- RISC + x86: ??%/year 2002 to present

Limit #3: Chip Yield


Manufacturing costs and yield problems limit use of density



- Moore's (Rock's) 2nd law: fabrication costs go up
- Yield (% usable chips) drops
- Parallelism can help
 - More smaller, simpler processors are easier to design and validate
 - Can use partially working chips:
 - E.g., Cell processor (PS3) is sold with 7 out of 8 “on” to improve yield

Revolution is Happening Now

- Chip density is continuing increase ~2x every 2 years
 - Clock speed is not
 - Number of processor cores may double instead
- There is little or no hidden parallelism (ILP) to be found
- Parallelism must be exposed to and managed by software



Source: Intel, Microsoft (Sutter) and Stanford (Olukotun, Hammond)

Multicore in Products

- “We are dedicating all of our future product development to multicore designs. ... This is a sea change in computing”


Paul Otellini, President, Intel (2005)

- All microprocessor companies switch to MP (2X CPUs / 2 yrs)


Manufacturer/Year	AMD/'12	Intel/'13	IBM/'13	Sun/'07
Processors/chip	16	8	12	8
Threads/Processor	one time, 1PPA has 256 cores, for 7 Watts	16	8	16
Threads/chip cores	16	32	96	128

- The Intel Xeon Phi 7100 has 61 cores

Single-core computer



Single-Core CPU Chip




Intel Nehalem Register File

- **16 64-bit general purpose registers**
- **6 16-bit segment register**
- **64-bit instruction pointer register**
- **64-bit RFLAGS register**
- **8 64-bit MMX registers**
- **16 128-bit XMM registers**
- **32-bit MXCSR register**
- **60 Performance Registers**

Multi-Core Architectures

- This lecture is about a new trend in computer architecture:
Replicate multiple processor cores on a single die.



Multi-core CPU chip

Performance Registers (1)


- Level 1 data cache misses**
- Level 1 instruction cache misses**
- Level 2 data cache misses**
- Level 2 instruction cache misses**
- Level 3 data cache misses**
- Level 3 instruction cache misses**
- Level 1 cache misses**
- Level 2 cache misses**
- Level 3 cache misses**
- Requests for a snoop**
- Requests for exclusive access to shared cache line**
- Requests for exclusive access to clean cache line**
- Requests for cache line invalidation**
- Requests for cache line intervention**
- Level 3 load misses**
- Level 3 store misses**
- Cycles branch units are idle**
- Cycles integer units are idle**
- Cycles floating point units are idle**
- Cycles load/store units are idle**
- Data translation lookaside buffer misses**
- Instruction translation lookaside buffer misses**
- Total translation lookaside buffer misses**

Performance Registers (2)


Level 1 load misses
Level 1 store misses
Level 2 load misses
Level 2 store misses
Branch target address cache misses
Data prefetch cache misses
Level 3 data cache hits
Translation lookaside buffer shootdowns
Failed store conditional instructions
Successful store conditional instructions
Total store conditional instructions
Cycles Stalled Waiting for memory accesses
Cycles Stalled Waiting for memory Reads
Cycles Stalled Waiting for memory writes
Cycles with no instruction issue
Cycles with maximum instruction issue
Cycles with no instructions completed
Cycles with maximum instructions completed
Hardware interrupts
Floating point operations; optimized to count scaled single precision vector operations
Floating point operations; optimized to count scaled double precision vector operations
Single precision vector/SIMD instructions
Double precision vector/SIMD instructions

Multi-Core CPU chip


- The cores fit on a single processor socket
- Also called CMP (Chip Multi-Processor)



The cores run in parallel



Within each core, threads are time-sliced (just like on a uniprocessor)




Interactions with the Operating System

- OS perceives each core as a separate processor
- OS scheduler maps threads/processes to different cores
- Most major OS support multi-core today: Windows, Linux, Mac OS X, ...


Memory Hierarchy

Key Observations


- Access to L1 cache is on order of 1 cycle
- Access to L2 on order of 1 to 10 cycles
- Access to Main memory ~ 100's cycles
- Access to Disk ~ 1000's cycles




Processor and Memory are Far Apart




Reading from Memory




Reading from Memory




Reading from Memory




Writing to Memory




Writing to Memory




Writing to Memory




Cache: Reading from Memory




Cache: Reading from Memory




Cache: Reading from Memory




Cache Hit




Cache Hit




Cache Miss



Cache Miss



Cache Miss



Instruction Level Parallelism (ILP)

- **Parallelism at the machine-instruction level**
- The processor can re-order, pipeline instructions, split them into microinstructions, do aggressive branch prediction, etc.
- **Instruction-level parallelism enabled rapid increases in processor speeds over the last 15 years**

Instruction Level Parallelism (ILP)

Programme séquentiel

N'y aurait-il pas des instructions indépendante qui pourraient être exécutées en parallèle?

Comment générer de l'ILP?

Pipe-line du processeur

- Recouvrement d'exécution d'instructions

- Limité par la divisibilité de l'instruction

Superscalaire

- Plusieurs unités fonctionnelles

- Limité par le parallélisme intrinsèque du programme seq.

Comment accroître l'ILP?


Prédiction sur les branchements conditionnels

Réordonner les instructions (Out of Order Execution)


Recherche

Domaines de l'Architecture-compilation

Simultaneous Multithreading (SMT), Hyperthreading



Scalar Architecture



Reduced Instruction Set (RISC) Architecture:

Les instructions load/store font référence à la mémoire


Les unités fonctionnelles travaillent sur des données stockées dans les registres

Hiérarchie mémoire dans une architecture scalaire :

Les éléments utilisés récemment sont copiés dans le **cache**,

Les accès au cache sont plus rapides que les accès à la mémoire.

Scalar Processor



Réduire le temps des changements de contexte (cases grisées-bleues)

Accroître l'utilisation des unités d'exécution (cases blanches)

Lightweight Process/Thread

Objectifs

Mener plusieurs activités indépendantes au sein d'un processus

Exploitation des architectures SMP

Améliorer l'utilisation du processeurs (context-switch)

Exemples

Simulations

Serveurs de fichiers

Systèmes d'exploitation (!)


Solution sans l'aide du multithreading

Automate à états finis implanté « à la main »
(sauvegardes d'états)

Lightweight Process/Thread (2)

Principe

Détacher flot d'exécution et ressources



Introduits dans divers langages & systèmes


Programmation concurrente

Recouvrement des E/S

Exploitation des architectures SMP

Thread and Superscalar Processor

Processeur SuperScalaire



SMT, Hyperthreading

Evolution des Architectures

Phase 1 : Processeurs Multithreadés

Modification de l'architecture du processeur

Incorporer au processeur deux (ou plus) jeux de registres pour les contextes des threads

- Registres généraux

- Program Counter (PC), registre d'instruction

- Process Status Word (PSW), registre d'état

A tout instant, un thread et son contexte sont actifs

Changement du contexte courant instantané

- Appel système, IT

- Défaut de cache

Processeur IBM PowerPC RS 64

Recherche, non commercialisé

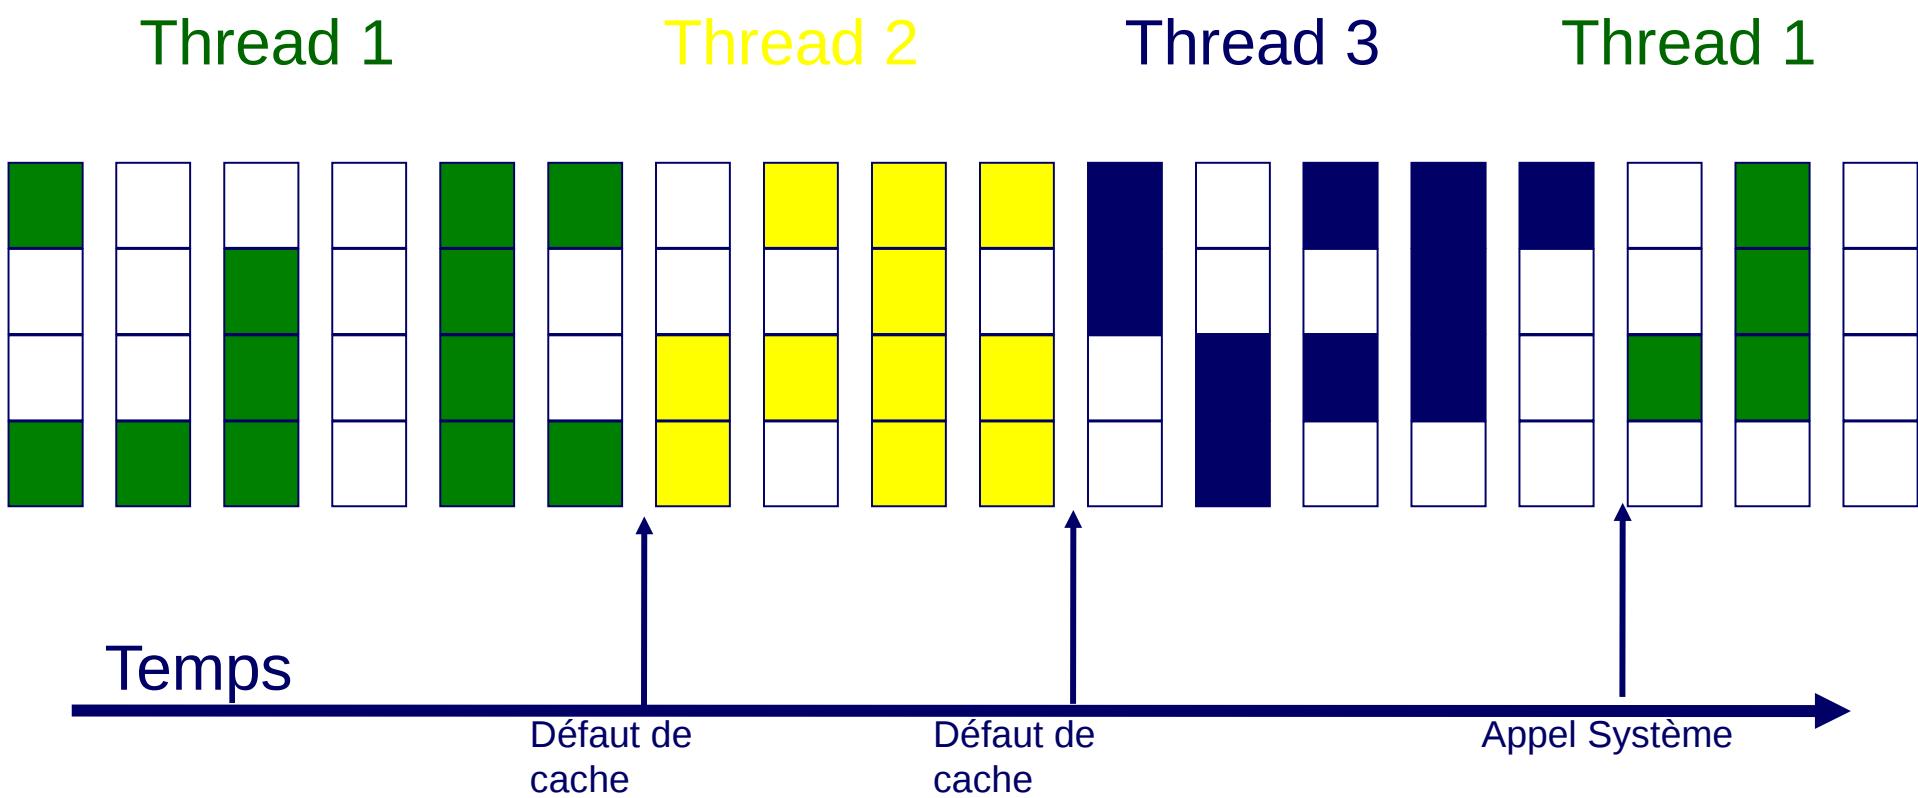
Processeur Intel Xeon Hyperthreading

Serveur Bi-processeur Xeon Nehalem Hyperthreadé

- Vue de Linux ou Windows, 4 processeurs

Phase 1 : Processeurs Multithreadés (2)

Multithread à Gros Grain (Coarse-Grained Multi-threaded)



Phase 2 : Processeurs Multithreadés

Modification de l'architecture du processeur

Incorporer au processeur **N** jeux de registres pour les contextes des threads

- Registres généraux

- Program Counter (PC), registre d'instruction

- Process Status Word (PSW), registre d'état

A tout instant, un thread et son contexte sont actifs

Changement de contexte à chaque cycle


- Chaque thread dispose de son ratio du processeur ($1/N$)

Processeur TERA

Phase 2 : Processeurs Multithreadés (2)

Multithread à Grain Fin (Fine-Grained Multi-threaded)

4 registres de threads : $\frac{1}{4}$ temps processeur par thread



Thread 1 Thread 2 Thread 3 Thread 4

Phase 3 : Processeurs Multithreadés

Modification de l'architecture du processeur

Incorporer au processeur **N** jeux de registres pour les contextes des threads

- Registres généraux

- Program Counter (PC), registre d'instruction

- Process Status Word (PSW), registre d'état


A un instant donné, les unités du processeurs peuvent être partagées entre plusieurs threads

Eviter de stresser les mêmes ressources


Conflit d'accès à certaines unités d'exécution

Processeurs ALPHA EV8, CRAY-TERA

Phase 3 : Processeurs Multithreadés (3)




Processeur Cray-TERA




Simultaneous Multithreading (SMT)

- Permits multiple independent threads to execute **SIMULTANEOUSLY** on the **SAME** core
- Weaving together multiple “threads” on the same core:
 - Exemple: If one thread is waiting for a floating point operation to complete, another thread can use the integer units

Without SMT, only a single thread can run at any given time




Without SMT, only a single thread can run at any given time



Thread 2:
integer operation

SMT processor: both threads can run concurrently




Thread 2: integer operation Thread 1: floating point operation


SMT not a « true » parallel processor

- Enables better threading (e.g. up to 30%)
- OS and applications perceives each simultaneous thread as a separate “virtual processor”
- The chip has only a single copy of each resource (functional units, pipeline,...)
- Compare to multi-core: each core has its own copy of resources

Multi-Core : threads can run on separate cores




Multi-Core : threads can run on separate cores



Combining Multi-core and SMT

- Cores can be SMT-enabled (or not)
- The different combinations
 - Single-core, non-SMT (standard uniprocessor)
 - Single-core, with SMT
 - Multi-core, non-SMT (ARM, AMD,...)
 - Multi-core, with SMT (Intel, IBM,...)
- The number of SMT threads: 2, 4 or sometimes 8 Simultaneous Threads
- Intel called them “Hyper Threads” (HT technology)

SMT Dual-core : all four threads can run concurrently



Comparison : multi-core vs SMT

- **Multi-core**
 - Since there are core, each is smaller and not as powerful (but also easier to design and manufacture)
 - However, great with thread-level parallelism
- **SMT**
 - Can have one large and fast superscalar core
 - Great performance on a single thread
 - Mostly still only exploit instruction level parallelism

The memory hierarchy for threading

- If simultaneous multithreading only:
 - All caches shared
- Multi-core chips
 - L1 private caches
 - L2 private caches in some architectures and shared in others
 - L3 shared caches (Last Level Cache on Intel)
- Memory is always shared


Private vs shared caches

- **Advantages of private caches**
 - They are closer to core, so faster access
 - No contention
- **Advantages of shared caches**
 - Threads on different cores can share the same cache data
 - More cache space available if a single (or few) intensive computing thread runs on the system

The cache coherence problem


- Since we have private caches
 - How to keep the data consistent across the caches
- Each core should perceive the memory as a monolithic array, shared by all cores

MESI
cache
Coheren
ce
Protocol



False sharing

- Performance issue in programs where core may write to different memory addresses BUT in the same cache lines
 - Known as Ping-Ponging – Cache line is shipped between cores



- False sharing is not an issue for shared caches
- It is an issue in private cache

Why Parallelism ?

- These arguments are no long theoretical
- All major processor vendors are producing multicore chips
 - Every machine is a parallel machine
 - All programmers will be parallel programmers???
- New software model
 - Want a new feature? Hide the “cost” by speeding up the code first
 - All programmers will be performance programmers???
- Some may eventually be hidden in libraries, compilers, and high level languages
 - But a lot of work is needed to get there
- Big open questions:
 - What will be the killer apps for multicore machines
 - How should the chips be designed, and how will they be programmed?


Outline

- Why powerful computers must be parallel processors
- Why writing (fast) parallel programs is hard
 - Including your laptop
- Principles of parallel computing performance
- Structure of the course

Why writing (fast) parallel programs is hard

Principles of Parallel Computing

- Finding enough parallelism (Amdahl's Law)
- Granularity
- Locality
- Load balance
- Coordination and synchronization
- Performance modeling



All of these things makes parallel programming even harder than sequential programming.

Finding Enough Parallelism

- Suppose only part of an application seems parallel
- Amdahl's law
 - let s be the fraction of work done sequentially, so $(1-s)$ is fraction parallelizable
 - $P = \text{number of processors}$

$$\text{Speedup}(P) = \text{Time}(1)/\text{Time}(P)$$

$$\leq 1/(s + (1-s)/P)$$

$$\leq 1/s$$



- Even if the parallel part speeds up perfectly performance is limited by the sequential part

Overhead of Parallelism

- Given enough parallel work, this is the biggest barrier to getting desired speedup
- Parallelism overheads include:
 - cost of starting a thread or process
 - cost of communicating shared data
 - cost of synchronizing
 - extra (redundant) computation
- Each of these can be in the range of milliseconds (=millions of flops) on some systems
- Tradeoff: Algorithm needs sufficiently large units of work to run fast in parallel (I.e. large granularity), but not so large that there is not enough parallel work

Locality and Parallelism

Conventional
Storage
Hierarchy



- Large memories are slow, fast memories are small
- Storage hierarchies are large and fast on average
- Parallel processors, collectively, have large, fast cache
 - the slow accesses to “remote” data we call “communication”
- Algorithm should do most work on local data

Load Imbalance

- Load imbalance is the time that some processors in the system are idle due to
 - insufficient parallelism (during that phase)
 - unequal size tasks
- Examples of the latter
 - adapting to “interesting parts of a domain”
 - tree-structured computations
 - fundamentally unstructured problems
- Algorithm needs to balance load

Reading Materials

- **Programming Many-Core Chips, Andras Vajda (Ericson), Springer, 2011**
- Ian Foster's book, “Designing and Building Parallel Programming”.
 - <http://www-unix.mcs.anl.gov/dbpp/>